

Problem 4: BloodMNIST (Data Loading and Exploration)

Q4.1: Data Loading and Exploration (2 points)

Tasks:

1. Install and load BloodMNIST using the `medmnist` package.
2. Report the number of samples in train/val/test splits.
3. Display a grid of 16 random training images (2 per class) with their class labels.
4. Plot the class distribution (bar chart) for the training set and comment on whether the dataset is balanced.

```
In [7]: # --- Task 1: Install and load BloodMNIST using the medmnist package ---
```

```
import os
import numpy as np
import matplotlib.pyplot as plt
from medmnist import INFO
from torchvision import transforms

data_dir = "./data"
os.makedirs(data_dir, exist_ok=True)

info = INFO["bloodmnist"]
BloodMNIST = getattr(medmnist, info["python_class"])
transform = transforms.ToTensor()

train_dataset = BloodMNIST(root=data_dir, split="train", transform=transform)
val_dataset = BloodMNIST(root=data_dir, split="val", transform=transform)
test_dataset = BloodMNIST(root=data_dir, split="test", transform=transform)

num_classes = len(info["label"])
class_names = [info["label"][str(i)] for i in range(num_classes)]
```

```
In [8]: # --- Task 2: Report the number of samples in train/val/test splits ---
```


```
n_train = len(train_dataset)
n_val = len(val_dataset)
n_test = len(test_dataset)
print(f"Training samples: {n_train}")
print(f"Validation samples: {n_val}")
print(f"Test samples: {n_test}")
```

```
Training samples: 11959
Validation samples: 1712
Test samples: 3421
```

```
In [11]: # --- Task 3: Display a grid of 16 random training images (2 per class) with
indices_per_class = {i: [] for i in range(num_classes)}
for idx, (img, target) in enumerate(train_dataset):
    y = int(target.squeeze()) if hasattr(target, "squeeze") else target
    if len(indices_per_class[y]) < 2:
        indices_per_class[y].append(idx)
    if all(len(v) == 2 for v in indices_per_class.values()):
        break

selected = [idx for c in range(num_classes) for idx in indices_per_class[c]]
fig, axes = plt.subplots(4, 4, figsize=(8, 8))
for ax, idx in zip(axes.flatten(), selected):
    img, target = train_dataset[idx]
    y = int(target.squeeze()) if hasattr(target, "squeeze") else target
    disp = img.permute(1, 2, 0).numpy() if img.dim() == 3 else img.squeeze()
    ax.imshow(disp, cmap="gray" if disp.ndim == 2 else None)
    title = class_names[y][:14] + (".." if len(class_names[y]) > 14 else "")
    ax.set_title(title, fontsize=9)
    ax.axis("off")
plt.suptitle("BloodMNIST: 16 training images (2 per class)")
plt.tight_layout()
plt.show()
```

BloodMNIST: 16 training images (2 per class)



```
In [12]: # --- Task 4: Plot class distribution (bar chart) for training set; comment
class_counts = np.zeros(num_classes, dtype=int)
for _, target in train_dataset:
    y = int(target.squeeze()) if hasattr(target, "squeeze") else target
    class_counts[y] += 1


plt.figure(figsize=(8, 4))
plt.bar(range(num_classes), class_counts)
# Shorten long class names for x-axis so the 4th class (e.g. immature granulocyte) fits
short_names = [s[:12] + "..." if len(s) > 12 else s for s in class_names]
plt.xticks(range(num_classes), short_names, rotation=45, ha="right", fontsize=10)
plt.ylabel("Count")
plt.title("BloodMNIST training set – class distribution")
plt.tight_layout()
plt.show()

# Distribution analysis
print("Class distribution (train):")
```

```

for i in range(num_classes):
    print(f" {class_names[i]:25s}: {class_counts[i]:4d} ({100*class_counts[i]:4.1f}%)")
print(f"\nMax: {class_counts.max()}\n{class_names[class_counts.argmax()]},\n"
print("The dataset is not balanced: some classes have many more samples than others")

```



Class distribution (train):

basophil	:	852 (7.1%)
eosinophil	:	2181 (18.2%)
erythroblast	:	1085 (9.1%)
immature granulocytes (myelocytes, metamyelocytes and promyelocytes)	:	2026 (16.9%)
lymphocyte	:	849 (7.1%)
monocyte	:	993 (8.3%)
neutrophil	:	2330 (19.5%)
platelet	:	1643 (13.7%)

Max: 2330 (neutrophil), Min: 849 (lymphocyte)

The dataset is not balanced: some classes have many more samples than others (e.g. neutrophil vs basophil).

Q4.2: Build and Train an MLP (5 points)

1. Flatten each 28×28×3 RGB image into a **2,352-dimensional** input vector
2. Build a 3-layer MLP: **Input 2352 → Hidden 1: 256 (ReLU) → Hidden 2: 128 (ReLU) → Output: 8**
3. **Adam optimizer, lr=1e-3, CrossEntropyLoss**
4. Train for **30 epochs**, batch size **64**
5. Plot **training loss** and **validation loss** (same figure)
6. Plot **training accuracy** and **validation accuracy** (same figure)

In [13]: # --- Q4.2: DataLoaders and device ---

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader

```

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
BATCH_SIZE = 64
train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffle=False)
```

In [14]: # --- Q4.2: 3-layer MLP (2352 → 256 ReLU → 128 ReLU → 8) ---

```
class BloodMLP(nn.Module):
    def __init__(self):
        super(BloodMLP, self).__init__()
        self.flatten = nn.Flatten()
        self.fc1 = nn.Linear(28 * 28 * 3, 256)
        self.fc2 = nn.Linear(256, 128)
        self.fc3 = nn.Linear(128, 8)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.flatten(x)
        x = self.relu(self.fc1(x))
        x = self.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

In [15]: # --- Q4.2: Model, optimizer, loss ---

```
model = BloodMLP().to(device)
optimizer = optim.Adam(model.parameters(), lr=1e-3)
criterion = nn.CrossEntropyLoss()
```

In [16]: # --- Q4.2: Train for 30 epochs, record train/val loss and accuracy ---

```
num_epochs = 30
train_losses, val_losses = [], []
train_accs, val_accs = [], []

for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    correct, total = 0, 0
    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)
        labels = labels.squeeze().long()
        optimizer.zero_grad()
        logits = model(images)
        loss = criterion(logits, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()
        correct += (logits.argmax(dim=1) == labels).sum().item()
        total += labels.size(0)

    train_losses.append(running_loss / len(train_loader))
    train_accs.append(correct / total)

    model.eval()
    val_loss = 0.0
    val_correct, val_total = 0, 0
```

```

with torch.no_grad():
    for images, labels in val_loader:
        images, labels = images.to(device), labels.to(device)
        labels = labels.squeeze().long()
        logits = model(images)
        val_loss += criterion(logits, labels).item()
        val_correct += (logits.argmax(dim=1) == labels).sum().item()
        val_total += labels.size(0)

    val_losses.append(val_loss / len(val_loader))
    val_accs.append(val_correct / val_total)
    print(f"Epoch {epoch+1}/{num_epochs} Train Loss: {train_losses[-1]:.4f}")

```

Epoch 1/30	Train Loss: 1.2534	Acc: 0.5322		Val Loss: 0.9025	Acc: 0.6583
Epoch 2/30	Train Loss: 0.8898	Acc: 0.6747		Val Loss: 0.7746	Acc: 0.7109
Epoch 3/30	Train Loss: 0.8050	Acc: 0.7062		Val Loss: 0.8001	Acc: 0.6805
Epoch 4/30	Train Loss: 0.7621	Acc: 0.7164		Val Loss: 0.6630	Acc: 0.7576
Epoch 5/30	Train Loss: 0.6974	Acc: 0.7444		Val Loss: 0.6148	Acc: 0.7862
Epoch 6/30	Train Loss: 0.6763	Acc: 0.7482		Val Loss: 0.6244	Acc: 0.7710
Epoch 7/30	Train Loss: 0.6333	Acc: 0.7669		Val Loss: 0.6758	Acc: 0.7494
Epoch 8/30	Train Loss: 0.5980	Acc: 0.7836		Val Loss: 0.5456	Acc: 0.7856
Epoch 9/30	Train Loss: 0.5835	Acc: 0.7849		Val Loss: 0.5466	Acc: 0.7996
Epoch 10/30	Train Loss: 0.5648	Acc: 0.7930		Val Loss: 0.5073	Acc: 0.8067
Epoch 11/30	Train Loss: 0.5337	Acc: 0.8042		Val Loss: 0.5134	Acc: 0.8172
Epoch 12/30	Train Loss: 0.5503	Acc: 0.7965		Val Loss: 0.6021	Acc: 0.7862
Epoch 13/30	Train Loss: 0.5283	Acc: 0.8057		Val Loss: 0.5374	Acc: 0.8008
Epoch 14/30	Train Loss: 0.5407	Acc: 0.7979		Val Loss: 0.5125	Acc: 0.8166
Epoch 15/30	Train Loss: 0.5125	Acc: 0.8113		Val Loss: 0.4635	Acc: 0.8370
Epoch 16/30	Train Loss: 0.4931	Acc: 0.8215		Val Loss: 0.5035	Acc: 0.8072
Epoch 17/30	Train Loss: 0.5013	Acc: 0.8146		Val Loss: 0.5140	Acc: 0.8032
Epoch 18/30	Train Loss: 0.4962	Acc: 0.8158		Val Loss: 0.4541	Acc: 0.8294
Epoch 19/30	Train Loss: 0.5113	Acc: 0.8109		Val Loss: 0.4743	Acc: 0.8248
Epoch 20/30	Train Loss: 0.4685	Acc: 0.8309		Val Loss: 0.5546	Acc: 0.8032
Epoch 21/30	Train Loss: 0.4682	Acc: 0.8292		Val Loss: 0.5271	Acc: 0.8172
Epoch 22/30	Train Loss: 0.4521	Acc: 0.8328		Val Loss: 0.5013	Acc: 0.8107
Epoch 23/30	Train Loss: 0.4582	Acc: 0.8292		Val Loss: 0.4943	Acc: 0.8148
Epoch 24/30	Train Loss: 0.4594	Acc: 0.8305		Val Loss: 0.5652	Acc: 0.7932
Epoch 25/30	Train Loss: 0.4482	Acc: 0.8368		Val Loss: 0.4837	Acc: 0.8143
Epoch 26/30	Train Loss: 0.4328	Acc: 0.8401		Val Loss: 0.5135	Acc: 0.8201
Epoch 27/30	Train Loss: 0.4313	Acc: 0.8401		Val Loss: 0.4533	Acc: 0.8376
Epoch 28/30	Train Loss: 0.4319	Acc: 0.8407		Val Loss: 0.5386	Acc: 0.8032
Epoch 29/30	Train Loss: 0.4354	Acc: 0.8428		Val Loss: 0.4598	Acc: 0.8283
Epoch 30/30	Train Loss: 0.4322	Acc: 0.8415		Val Loss: 0.4325	Acc: 0.8411

In [17]: # --- Q4.2: Plot training vs validation loss and accuracy ---

```

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(10, 4))
epochs = range(1, num_epochs + 1)
ax1.plot(epochs, train_losses, label="Train loss")
ax1.plot(epochs, val_losses, label="Val loss")
ax1.set_xlabel("Epoch")
ax1.set_ylabel("Loss")
ax1.set_title("Training vs Validation Loss")
ax1.legend()


ax2.plot(epochs, train_accs, label="Train acc")
ax2.plot(epochs, val_accs, label="Val acc")

```

```

ax2.set_xlabel("Epoch")
ax2.set_ylabel("Accuracy")
ax2.set_title("Training vs Validation Accuracy")
ax2.legend()
plt.tight_layout()
plt.show()

```



Q4.3: Evaluation and Analysis (5 points)

1. Report the **final test accuracy** (as a percentage)
2. Generate a **confusion matrix** (8×8) for the test set using
 `sklearn.metrics.confusion_matrix` and visualize it as a **heatmap** with class labels
3. Identify the **two most confused cell type pairs** (highest off-diagonal values)
4. Compute and report **per-class precision and recall** using
 `sklearn.metrics.classification_report`
5. Which cell type has the **lowest recall**? Examine **5 misclassified examples** of this cell type and hypothesize why the model struggles with it.

```

In [18]: # --- Q4.3.1: Final test accuracy ---
from torch.utils.data import DataLoader
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)

model.eval()
all_preds, all_labels = [], []
with torch.no_grad():
    for images, labels in test_loader:
        images = images.to(device)
        labels = labels.squeeze().long()
        logits = model(images)
        all_preds.append(logits.argmax(dim=1).cpu())
        all_labels.append(labels.cpu())

all_preds = torch.cat(all_preds).numpy()
all_labels = torch.cat(all_labels).numpy()

```

```

test_acc = 100.0 * (all_preds == all_labels).mean()
print(f"Final test accuracy: {test_acc:.2f}%")

```


Final test accuracy: 82.96%

In [20]: # --- Q4.3.2: Confusion matrix (8x8) heatmap with class labels ---
from sklearn.metrics import confusion_matrix

```

cm = confusion_matrix(all_labels, all_preds)
fig, ax = plt.subplots(figsize=(8, 6))
im = ax.imshow(cm, cmap="Blues")
plt.colorbar(im, ax=ax, label="Count")
ax.set_xticks(range(8))
ax.set_yticks(range(8))
short_names = [s[:12] + '...' if len(s) > 12 else s for s in class_names]
ax.set_xticklabels(short_names, rotation=45, ha="right", fontsize=9)
ax.set_yticklabels(short_names, fontsize=9)
ax.set_xlabel("Predicted")
ax.set_ylabel("True")
for i in range(8):
    for j in range(8):
        ax.text(j, i, cm[i, j], ha="center", va="center", color="black" if cm[i, j] > 100 else "white")
plt.title("Confusion matrix (test set)")
plt.tight_layout()
plt.show()

```



```
In [21]: # --- Q4.3.3: Two most confused cell type pairs (highest off-diagonal) ---
# Copy cm and zero out diagonal to get only off-diagonal
cm_off = cm.copy()
for i in range(8):
    cm_off[i, i] = 0

# Find indices of 2 largest off-diagonal values (flat index)
flat = cm_off.flatten()
top2_flat = np.argsort(flat)[-2:][::-1]
for idx in top2_flat:
    i, j = idx // 8, idx % 8
    print(f" {class_names[i]} predicted as {class_names[j]}: {cm[i, j]} times")
print("Two most confused pairs above.)")
```

immature granulocytes(myelocytes, metamyelocytes and promyelocytes) predicted as monocyte: 83 times
 immature granulocytes(myelocytes, metamyelocytes and promyelocytes) predicted as neutrophil: 62 times
 (Two most confused pairs above.)

```
In [24]: # --- Q4.3.4: Per-class precision and recall (classification_report) ---
from sklearn.metrics import classification_report

print(classification_report(all_labels, all_preds, target_names=class_names,
```

				precision	
	on	recall	f1-score	support	
62	basophil	0.652	0.657	244	0.6
84	eosinophil	0.976	0.928	624	0.8
19	erythroblast	0.801	0.856	311	0.9
	immature granulocytes(myelocytes, metamyelocytes and promyelocytes)				0.7
36	lymphocyte	0.554	0.633	579	0.8
08	monocyte	0.778	0.792	243	0.6
34	neutrophil	0.757	0.690	284	0.8
59	platelet	0.943	0.899	666	0.9
73		0.996	0.984	470	
				accuracy	
0.830		3421			
09	macro avg	0.807	0.805	3421	0.8
28	weighted avg	0.830	0.825	3421	0.8

```
In [25]: # --- Q4.3.5: Cell type with lowest recall; show 5 misclassified examples ---
from sklearn.metrics import precision_recall_fscore_support
precision, recall, _, _ = precision_recall_fscore_support(all_labels, all_pr
```

```


lowest_recall_class = int(np.argmin(recall))
print(f"Cell type with lowest recall: {class_names[lowest_recall_class]} (recall = {recall[lowest_recall_class]:.4f})")

# Indices in test set where true=lowest_recall_class but predicted != lowest_recall_class
mis_mask = (all_labels == lowest_recall_class) & (all_preds != lowest_recall_class)
mis_indices = np.where(mis_mask)[0][:5]

short_names = [s[:12] + '...' if len(s) > 12 else s for s in class_names]
fig, axes = plt.subplots(1, 5, figsize=(12, 3))
for ax, idx in zip(axes, mis_indices):
    img, _ = test_dataset[idx]
    true_lab = all_labels[idx]
    pred_lab = all_preds[idx]
    disp = img.permute(1, 2, 0).numpy() if img.dim() == 3 else img.squeeze()
    ax.imshow(disp, cmap="gray" if disp.ndim == 2 else None)
    ax.set_title(f"True: {short_names[true_lab]}\nPred: {short_names[pred_lab]}")
    ax.axis("off")
plt.suptitle(f"5 misclassified examples (true label = {short_names[lowest_recall_class]})")
plt.tight_layout()
plt.show()

```

Cell type with lowest recall: immature granulocytes (myelocytes, metamyelocytes and promyelocytes) (recall = 0.554)



Hypothesis (why the model struggles with this cell type):

The class with lowest recall is often a minority class or one that is visually similar to others (e.g. immature granulocytes vs neutrophils). The model may confuse it with the most confused pair identified above, or the training set may have few examples. Inspect the 5 misclassified images: shared appearance (size, shape, color) with the predicted class can explain the errors.

Q4.4: Prediction Confidence Analysis (3 points)

Categorize predictions into four quadrants by **confidence** (max softmax prob) and **correctness**:

- **High confidence:** max prob > 0.9
- **Low confidence:** max prob < 0.6

	Correct	Incorrect
High confidence	✓ Confident	✗ Confident
Low confidence	✓ Uncertain	✗ Uncertain

1. Find and display **2 examples from each quadrant** (8 images total) with: image, true label, predicted label, prediction confidence.
2. **Written analysis (3–4 sentences):** What distinguishes "Incorrect but Confident" examples? Why might the model be overconfident?

```
In [26]: # --- Q4.4: Get max softmax probability for each test sample ---
model.eval()
all_probs = []
with torch.no_grad():
    for images, _ in test_loader:
        images = images.to(device)
        logits = model(images)
        probs = torch.softmax(logits, dim=1)
        max_prob, _ = probs.max(dim=1)
        all_probs.append(max_prob.cpu().numpy())
all_probs = np.concatenate(all_probs)

correct = (all_preds == all_labels)
high_conf = (all_probs > 0.9)
low_conf = (all_probs < 0.6)

# Four quadrants: (correct, high), (correct, low), (incorrect, high), (incorrect, low)
quadrants = [
    ("Correct, High conf", np.where(correct & high_conf)[0]),
    ("Correct, Low conf", np.where(correct & low_conf)[0]),
    ("Incorrect, High conf", np.where(~correct & high_conf)[0]),
    ("Incorrect, Low conf", np.where(~correct & low_conf)[0]),
]
for name, idx in quadrants:
    print(f" {name}: {len(idx)} samples")

Correct, High conf: 1927 samples
Correct, Low conf: 287 samples
Incorrect, High conf: 53 samples
Incorrect, Low conf: 302 samples
```

```
In [27]: # --- Q4.4: Display 2 examples from each quadrant (8 images total) ---
short_names = [s[:12] + ".." if len(s) > 12 else s for s in class_names]
fig, axes = plt.subplots(4, 2, figsize=(8, 12))
for row, (quad_name, indices) in enumerate(quadrants):
    chosen = indices[2:]
    for col, idx in enumerate(chosen):
        ax = axes[row, col]
        img, _ = test_dataset[idx]
        disp = img.permute(1, 2, 0).numpy() if img.dim() == 3 else img.squeeze(0)
        ax.imshow(disp, cmap="gray" if disp.ndim == 2 else None)
        true_lab = all_labels[idx]
        pred_lab = all_preds[idx]
        conf = all_probs[idx]
        ax.set_title(f"True: {short_names[true_lab]}\nPred: {short_names[pred_lab]}")
        ax.axis("off")
    axes[row, 0].set_ylabel(quad_name, fontsize=10)
plt.suptitle("Q4.4: 2 examples per quadrant (Correct/Incorrect x High/Low confidence)")
```

```
plt.tight_layout()  
plt.show()
```

Q4.4: 2 examples per quadrant (Correct/Incorrect × High/Low confidence)

True: basophil
Pred: basophil
Conf: 0.945



True: eosinophil
Pred: eosinophil
Conf: 0.995




True: immature gra..
Pred: immature gra..
Conf: 0.459



True: monocyte
Pred: monocyte
Conf: 0.453



True: neutrophil
Pred: eosinophil
Conf: 0.939




True: neutrophil
Pred: erythroblast
Conf: 0.916




True: immature gra..
Pred: neutrophil
Conf: 0.455



True: eosinophil
Pred: immature gra..
Conf: 0.436





Written analysis (3–4 sentences):

The "Incorrect but Confident" examples are cases where the model assigns high probability (>0.9) to the wrong class. Often these are cell types that look visually similar to the predicted class (e.g. granulocytes vs neutrophils, or similar staining and shape), so the model has learned features that confidently but wrongly distinguish them. The model may be overconfident on these because the training data or the 28×28 resolution does not capture the subtle differences that a pathologist would use, so the classifier relies on strong but misleading cues and assigns high confidence to the wrong class.